

Simulation and High-Performance Computing

Part 5: Finite Difference Methods for Partial Differential Equations

Steffen Börm

Christian-Albrechts-Universität zu Kiel

September 30th, 2020

Poisson's equation

Model problem: Poisson's equation on the unit square

$$-\Delta u(x) = f(x) \quad \text{for all } x \in \Omega := (0, 1) \times (0, 1),$$

Laplace operator: Partial differential operator of second order,

$$\Delta u(x) = \frac{\partial^2 u}{\partial x_1^2}(x) + \frac{\partial^2 u}{\partial x_2^2}(x).$$

Poisson's equation

Model problem: Poisson's equation on the unit square

$$\begin{aligned} -\Delta u(x) &= f(x) && \text{for all } x \in \Omega := (0, 1) \times (0, 1), \\ u(x) &= 0 && \text{for all } x \in \partial\Omega = \{0, 1\} \times [0, 1] \cup [0, 1] \times \{0, 1\}, \end{aligned}$$

with Dirichlet boundary conditions.

Laplace operator: Partial differential operator of second order,

$$\Delta u(x) = \frac{\partial^2 u}{\partial x_1^2}(x) + \frac{\partial^2 u}{\partial x_2^2}(x).$$

Applications: Electrostatic fields, heat dissipation, wave propagation, fluid dynamics, ...

Grids and grid functions


Problem: The domain Ω contains infinitely many points, but a computer can only store finitely many values.

Grid: Replace the domain by a finite number of points.

Choose $N \in \mathbb{N}$, let $h := \frac{1}{N+1}$ and

$$\Omega_h := \{(ih, jh) : i, j \in [1 : N]\}, \quad \bar{\Omega}_h := \Omega_h \cup \partial\Omega_h,$$

$$\partial\Omega_h := \{(ih, jh) : i, j \in [0 : N+1], i \in \{0, N+1\} \vee j \in \{0, N+1\}\}.$$



Grids and grid functions



Problem: The domain Ω contains infinitely many points, but a computer can only store finitely many values.

Grid: Replace the domain by a finite number of points.

Choose $N \in \mathbb{N}$, let $h := \frac{1}{N+1}$ and

$$\Omega_h := \{(ih, jh) : i, j \in [1 : N]\}, \quad \bar{\Omega}_h := \Omega_h \cup \partial\Omega_h,$$

$$\partial\Omega_h := \{(ih, jh) : i, j \in [0 : N+1], i \in \{0, N+1\} \vee j \in \{0, N+1\}\}.$$



Grids and grid functions




Problem: The domain Ω contains infinitely many points, but a computer can only store finitely many values.

Grid: Replace the domain by a finite number of points.

Choose $N \in \mathbb{N}$, let $h := \frac{1}{N+1}$ and

$$\Omega_h := \{(ih, jh) : i, j \in [1 : N]\}, \quad \bar{\Omega}_h := \Omega_h \cup \partial\Omega_h,$$

$$\partial\Omega_h := \{(ih, jh) : i, j \in [0 : N+1], i \in \{0, N+1\} \vee j \in \{0, N+1\}\}.$$






Grids and grid functions

Problem: The domain Ω contains infinitely many points, but a computer can only store finitely many values.

Grid: Replace the domain by a finite number of points.

Choose $N \in \mathbb{N}$, let $h := \frac{1}{N+1}$ and

$$\begin{aligned}\Omega_h &:= \{(ih, jh) : i, j \in [1 : N]\}, & \bar{\Omega}_h &:= \Omega_h \cup \partial\Omega_h, \\ \partial\Omega_h &:= \{(ih, jh) : i, j \in [0 : N+1], i \in \{0, N+1\} \vee j \in \{0, N+1\}\}.\end{aligned}$$



Grid function: Replace a function $u : \bar{\Omega} \rightarrow \mathbb{R}$ by a mapping $u_h : \bar{\Omega}_h \rightarrow \mathbb{R}$.

Finite difference approximation

Problem: We cannot evaluate derivatives of a grid function, since it is only defined in discrete points.

Approach: Replace differential operators by difference quotients.

$$g''(t) \approx \frac{g(t+h) - 2g(t) + g(t-h)}{h^2},$$

Finite difference approximation

Problem: We cannot evaluate derivatives of a grid function, since it is only defined in discrete points.

Approach: Replace differential operators by difference quotients.

$$g''(t) \approx \frac{g(t+h) - 2g(t) + g(t-h)}{h^2},$$

$$\begin{aligned}\Delta u(x) &= \frac{\partial^2 u}{\partial x_1^2}(x) + \frac{\partial^2 u}{\partial x_2^2}(x) \\ &\approx \frac{u(x_1+h, x_2) - 2u(x) + u(x_1-h, x_2)}{h^2} \\ &\quad + \frac{u(x_1, x_2+h) - 2u(x) + u(x_1, x_2-h)}{h^2} =: \Delta_h u(x).\end{aligned}$$

Finite difference approximation

Problem: We cannot evaluate derivatives of a grid function, since it is only defined in discrete points.

Approach: Replace differential operators by difference quotients.

$$g''(t) \approx \frac{g(t+h) - 2g(t) + g(t-h)}{h^2},$$

$$\begin{aligned} \Delta u(x) &= \frac{\partial^2 u}{\partial x_1^2}(x) + \frac{\partial^2 u}{\partial x_2^2}(x) \\ &\approx \frac{u(x_1+h, x_2) - 2u(x) + u(x_1-h, x_2)}{h^2} \\ &\quad + \frac{u(x_1, x_2+h) - 2u(x) + u(x_1, x_2-h)}{h^2} =: \Delta_h u(x). \end{aligned}$$

Finite difference approximation

Problem: We cannot evaluate derivatives of a grid function, since it is only defined in discrete points.

Approach: Replace differential operators by difference quotients.

$$g''(t) \approx \frac{g(t+h) - 2g(t) + g(t-h)}{h^2},$$

$$\begin{aligned}\Delta u(x) &= \frac{\partial^2 u}{\partial x_1^2}(x) + \frac{\partial^2 u}{\partial x_2^2}(x) \\ &\approx \frac{u(x_1+h, x_2) - 2u(x) + u(x_1-h, x_2)}{h^2} \\ &\quad + \frac{u(x_1, x_2+h) - 2u(x) + u(x_1, x_2-h)}{h^2} =: \Delta_h u(x).\end{aligned}$$

Finite difference approximation

Problem: We cannot evaluate derivatives of a grid function, since it is only defined in discrete points.

Approach: Replace differential operators by difference quotients.


$$\begin{aligned} g''(t) &\approx \frac{g(t+h) - 2g(t) + g(t-h)}{h^2}, \\ \Delta u(x) &= \frac{\partial^2 u}{\partial x_1^2}(x) + \frac{\partial^2 u}{\partial x_2^2}(x) \\ &\approx \frac{u(x_1+h, x_2) - 2u(x) + u(x_1-h, x_2)}{h^2} \\ &\quad + \frac{u(x_1, x_2+h) - 2u(x) + u(x_1, x_2-h)}{h^2} =: \Delta_h u(x). \end{aligned}$$

Important: If $x \in \Omega_h$, the approximation $\Delta_h u_h(x)$ is well-defined.

Discrete Laplace operator

Five-point stencil: The Laplace operator is approximated by


$$\Delta_h u_h(x) = \frac{u(x_1 + h, x_2) + u(x_1 - h, x_2) + u(x_1, x_2 + h) + u(x_1, x_2 - h) - 4u(x)}{h^2}$$



Discrete Laplace operator

Five-point stencil: The Laplace operator is approximated by


$$\Delta_h u_h(x) = \frac{u(x_1 + h, x_2) + u(x_1 - h, x_2) + u(x_1, x_2 + h) + u(x_1, x_2 - h) - 4u(x)}{h^2}$$



Discrete Laplace operator

Five-point stencil: The Laplace operator is approximated by

$$\Delta_h u_h(x) = \frac{u(x_1 + h, x_2) + u(x_1 - h, x_2) + u(x_1, x_2 + h) + u(x_1, x_2 - h) - 4u(x)}{h^2}$$




Boundary points are a special case, since they are not degrees of freedom.

Discrete Laplace operator

Five-point stencil: The Laplace operator is approximated by

$$\Delta_h u_h(x) = \frac{u(x_1 + h, x_2) + u(x_1 - h, x_2) + u(x_1, x_2 + h) + u(x_1, x_2 - h) - 4u(x)}{h^2}$$




Boundary points are a special case, since they are not degrees of freedom.

Discrete Laplace operator

Five-point stencil: The Laplace operator is approximated by

$$\Delta_h u_h(x) = \frac{u(x_1 + h, x_2) + u(x_1 - h, x_2) + u(x_1, x_2 + h) + u(x_1, x_2 - h) - 4u(x)}{h^2}$$



Boundary points are a special case, since they are not degrees of freedom.

Linear system

Boundary value problem: Find $u \in C(\bar{\Omega})$ with $u|_{\Omega} \in C^2(\Omega)$ such that

$$\begin{aligned} -\Delta u(x) &= f(x) && \text{for all } x \in \Omega, \\ u(x) &= 0 && \text{for all } x \in \partial\Omega. \end{aligned}$$

Finite difference approximation: Find $u_h: \bar{\Omega}_h \rightarrow \mathbb{R}$ such that

$$\begin{aligned} -\Delta_h u_h(x) &= f(x) && \text{for all } x \in \Omega_h, \\ u_h(x) &= 0 && \text{for all } x \in \partial\Omega_h. \end{aligned}$$

Linear system

Boundary value problem: Find $u \in C(\bar{\Omega})$ with $u|_{\Omega} \in C^2(\Omega)$ such that

$$\begin{aligned} -\Delta u(x) &= f(x) && \text{for all } x \in \Omega, \\ u(x) &= 0 && \text{for all } x \in \partial\Omega. \end{aligned}$$

Finite difference approximation: Find $u_h: \bar{\Omega}_h \rightarrow \mathbb{R}$ such that

$$\begin{aligned} -\Delta_h u_h(x) &= f(x) && \text{for all } x \in \Omega_h, \\ u_h(x) &= 0 && \text{for all } x \in \partial\Omega_h. \end{aligned}$$

Result: We approximate the partial differential equation by a linear system of dimension $n := N^2$.

Lexicographic enumeration

Problem: The index set $\bar{\Omega}_h$ is two-dimensional, but standard programming languages prefer one-dimensional structures.

Approach: Enumerate row by row.

$$x = (ih, jh) \mapsto i + (N + 2)j \quad i, j \in [0 : N + 1].$$

Implementation: Grid functions represented by arrays with $(N + 2)^2$ elements, $u_h(ih, jh) = u[i + j * \text{yinc}]$ with $\text{yinc} = N + 2$.

Lexicographic enumeration

Problem: The index set $\bar{\Omega}_h$ is two-dimensional, but standard programming languages prefer one-dimensional structures.

Approach: Enumerate row by row.

$$x = (ih, jh) \mapsto i + (N + 2)j \quad i, j \in [0 : N + 1].$$

Implementation: Grid functions represented by arrays with $(N + 2)^2$ elements, $u_h(ih, jh) = u[i + j * \text{yinc}]$ with $\text{yinc} = N + 2$.

Boundary points can be eliminated, since they are not degrees of freedom.

$$x = (ih, jh) \mapsto (i - 1) + N(j - 1) \quad i, j \in [1 : N],$$

this leads to arrays with $n := N^2$ elements containing only unknown grid values $u_h(ih, jh) = u[(i-1) + (j-1) * \text{yinc}]$ with $\text{yinc} = N$.

Matrix structure

Approach: Degrees of freedom in the j -th row of a grid function u_h are collected in a vector $u^{(j)} \in \mathbb{R}^N$ with $u_i^{(j)} := u_h(ih, jh)$, $i, j \in [1 : N]$.

Linear system is now n -dimensional with $n = N^2$, given by

$$\frac{1}{h^2} \begin{pmatrix} T & -I & & \\ -I & \ddots & \ddots & \\ & \ddots & \ddots & -I \\ & & -I & T \end{pmatrix} \begin{pmatrix} u^{(1)} \\ u^{(2)} \\ \vdots \\ u^{(N)} \end{pmatrix} = \begin{pmatrix} f^{(1)} \\ f^{(2)} \\ \vdots \\ f^{(N)} \end{pmatrix}$$

with the tridiagonal matrix

$$T := \begin{pmatrix} 4 & -1 & & \\ -1 & \ddots & \ddots & \\ & \ddots & \ddots & -1 \\ & & -1 & 4 \end{pmatrix} \in \mathbb{R}^{N \times N}.$$

Direct solvers

Problem: We have to solve the linear system $Au = f$ with the matrix

$$A := \frac{1}{h^2} \begin{pmatrix} T & -I & & \\ -I & \ddots & \ddots & \\ & \ddots & \ddots & -I \\ & & -I & T \end{pmatrix}, \quad T := \begin{pmatrix} 4 & -1 & & \\ -1 & \ddots & \ddots & \\ & \ddots & \ddots & -1 \\ & & -1 & 4 \end{pmatrix}.$$

First approach: If n is not too large, we can use a direct solver, e.g., Gaussian elimination.

Direct solvers

Problem: We have to solve the linear system $Au = f$ with the matrix

$$A := \frac{1}{h^2} \begin{pmatrix} T & -I & & \\ -I & \ddots & \ddots & \\ & \ddots & \ddots & -I \\ & & -I & T \end{pmatrix}, \quad T := \begin{pmatrix} 4 & -1 & & \\ -1 & \ddots & \ddots & \\ & \ddots & \ddots & -1 \\ & & -1 & 4 \end{pmatrix}.$$

First approach: If n is not too large, we can use a direct solver, e.g., Gaussian elimination.

Refined approach: Use Krylov methods or multigrid iteration, these are topics in parts 8 and 9 of this course.

Example: LU factorization

Triangular matrices:

- $L \in \mathbb{R}^{n \times n}$ is **lower triangular** if $\ell_{ij} = 0$ for all $i < j$.
- $U \in \mathbb{R}^{n \times n}$ is **upper triangular** if $u_{ij} = 0$ for all $i > j$.

Triangular systems $Lx = b$ and $Ux = b$ can be solved by forward and backward substitution.

LU factorization: A matrix $A \in \mathbb{R}^{n \times n}$ is split into $A = LU$.

$$Ax = b \quad \Longleftrightarrow \quad Ly = b \text{ and } Ux = y.$$

Forward substitution

Goal: Solve $Lx = b$ with a lower triangular matrix L .

Approach: Split into submatrices and -vectors.

$$L = \begin{pmatrix} \ell_{11} & \\ L_{*1} & L_{**} \end{pmatrix}, \quad x = \begin{pmatrix} x_1 \\ x_* \end{pmatrix}, \quad b = \begin{pmatrix} b_1 \\ b_* \end{pmatrix}.$$

We obtain

$$\begin{pmatrix} \ell_{11}x_1 \\ L_{*1}x_1 + L_{**}x_* \end{pmatrix} = \begin{pmatrix} \ell_{11} & \\ L_{*1} & L_{**} \end{pmatrix} \begin{pmatrix} x_1 \\ x_* \end{pmatrix} = Lx = b = \begin{pmatrix} b_1 \\ b_* \end{pmatrix}$$

and can solve $\ell_{11}x_1 = b_1$ directly and

$$L_{**}x_* = b_* - L_{*1}x_1$$

by recursion, since L_{**} is again lower triangular.

Backward substitution

Goal: Solve $Ux = b$ with an upper triangular matrix U .

Approach: Split into submatrices and -vectors.

$$U = \begin{pmatrix} U_{**} & U_{*n} \\ & u_{nn} \end{pmatrix}, \quad x = \begin{pmatrix} x_* \\ x_n \end{pmatrix}, \quad b = \begin{pmatrix} b_* \\ b_n \end{pmatrix}.$$

We obtain

$$\begin{pmatrix} U_{**}x_* + U_{*n}x_n \\ u_{nn}x_n \end{pmatrix} = \begin{pmatrix} U_{**} & U_{*n} \\ & u_{nn} \end{pmatrix} \begin{pmatrix} x_* \\ x_n \end{pmatrix} = Ux = b = \begin{pmatrix} b_* \\ b_n \end{pmatrix}$$

and can solve $u_{nn}x_n = b_n$ directly and

$$U_{**}x_* = b_* - U_{*n}x_n$$

by recursion, since U_{**} is again upper triangular.

LU factorization

Goal: Given A , find L lower and U upper triangular with $A = LU$.

Approach: Split into submatrices.

$$A = \begin{pmatrix} a_{11} & A_{1*} \\ A_{*1} & A_{**} \end{pmatrix}, \quad L = \begin{pmatrix} \ell_{11} & \\ L_{*1} & L_{**} \end{pmatrix}, \quad U = \begin{pmatrix} u_{11} & U_{1*} \\ & U_{**} \end{pmatrix}.$$

We obtain

$$\begin{pmatrix} \ell_{11}u_{11} & \ell_{11}U_{1*} \\ L_{*1}u_{11} & L_{**}U_{**} + L_{*1}U_{1*} \end{pmatrix} = \begin{pmatrix} \ell_{11} & \\ L_{*1} & L_{**} \end{pmatrix} \begin{pmatrix} u_{11} & U_{1*} \\ & U_{**} \end{pmatrix} = \begin{pmatrix} a_{11} & A_{1*} \\ A_{*1} & A_{**} \end{pmatrix}$$

and can solve $\ell_{11}u_{11} = a_{11}$ directly (usually $\ell_{11} = 1$, $u_{11} = a_{11}$),
 $L_{*1}u_{11} = A_{*1}$ and $\ell_{11}U_{1*} = A_{1*}$ by scaling, and use recursion for

$$L_{**}U_{**} = A_{**} - L_{*1}U_{1*}.$$

Arrays and pointers in C

Arrays and pointers: In C, arrays and pointers are almost synonymous. If A is a pointer, A[k] means accessing the element that can be found k steps behind the one A points to.

Pointer arithmetic: We can move pointers by adding integers.

```
real A[] = { 1.0, 2.0, 3.0, 4.0, 5.0 };  
real *B, *C;
```

```
B = A + 2;  
printf("%f\n", B[1]);    /* Yields B[1] = A[2+1] = 4.0 */
```

```
C = B - 1;  
printf("%f\n", C[3]);    /* Yields C[3] = B[2] = A[4] = 5.0 */
```

Array representation of matrices and vectors

Vectors $x \in \mathbb{R}^n$ are represented by

- a pointer x to the first coefficient,
- an increment $incx$ that takes us to the next coefficient, and
- the dimension n .

We can find x_i in $x[(i-1)*incx]$, $i \in [1 : n]$.

Array representation of matrices and vectors

Vectors $x \in \mathbb{R}^n$ are represented by

- a pointer x to the first coefficient,
- an increment $incx$ that takes us to the next coefficient, and
- the dimension n .

We can find x_i in $x[(i-1)*incx]$, $i \in [1 : n]$.

Matrices $A \in \mathbb{R}^{m \times n}$ are represented by

- a pointer A to the first coefficient a_{11} ,
- a leading dimension ldA that takes us to the next column, and
- the dimensions $rowsA$ and $colsA$.

We can find a_{ij} in $A[(i-1)+(j-1)*ldA]$, $i \in [1 : m]$, $j \in [1 : n]$.

Array representation of matrices and vectors

Vectors $x \in \mathbb{R}^n$ are represented by

- a pointer x to the first coefficient,
- an increment $incx$ that takes us to the next coefficient, and
- the dimension n .

We can find x_i in $x[(i-1)*incx]$, $i \in [1 : n]$.

Matrices $A \in \mathbb{R}^{m \times n}$ are represented by

- a pointer A to the first coefficient a_{11} ,
- a leading dimension ldA that takes us to the next column, and
- the dimensions $rowsA$ and $colsA$.

We can find a_{ij} in $A[(i-1)+(j-1)*ldA]$, $i \in [1 : m]$, $j \in [1 : n]$.

Standard practice: Better to use indices $[0 : n - 1]$ instead of $[1 : n]$.

Submatrices and -vectors

Submatrix: If $A \in \mathbb{R}^{n \times m}$ is given,


```
B = A + 2 + 3 * 1dA;
```

```
1dB = 1dA;
```

```
rowsB = rowsA - 2;
```

```
colsB = colsA - 3;
```

defines the submatrix starting in the third row and the fourth column.



Submatrices and -vectors

Submatrix: If $A \in \mathbb{R}^{n \times m}$ is given,


$$B = A + 2 + 3 * \text{ldA};$$

$$\text{ldB} = \text{ldA};$$

$$\text{rowsB} = \text{rowsA} - 2;$$

$$\text{colsB} = \text{colsA} - 3;$$

defines the submatrix starting in the third row and the fourth column.



Subvectors: The second row and the third column are represented by

$$x = A + 1; \quad \text{incx} = \text{ldA}; \quad n = \text{colsA};$$

$$y = A + 2 * \text{ldA}; \quad \text{incy} = 1; \quad m = \text{rowsA};$$

BLAS

Basic Linear Algebra Subprograms can be used to take advantage of highly optimized implementations of basic operations on matrices and vectors.

BLAS Level 1 contains vector operations, for example

- scaling a vector, $x \leftarrow \alpha x$:

```
void scal(int n, real alpha,  
          real *x, int incx);
```

- adding two vectors, $y \leftarrow y + \alpha x$:

```
void axpy(int n, real alpha,  
          const real *x, int incx,  
          real *y, int incy);
```


- computing the Euclidean norm $\|x\|_2 = (x_1^2 + \dots + x_n^2)^{1/2}$:

```
real nrm2(int n, const real *x, int incx);
```

Forward substitution with BLAS

Goal: Solve $Lx = b$. Using the block notation

$$\begin{pmatrix} \ell_{11} & \\ L_{*1} & L_{**} \end{pmatrix} \begin{pmatrix} x_1 \\ x_* \end{pmatrix} = \begin{pmatrix} b_1 \\ b_* \end{pmatrix},$$



equivalent with $x_1 = b_1 / \ell_{11}$ and $L_{**}x_* = b_* - L_{*1}x_1$.


Recursion: We have to apply the algorithm to the submatrix L_{**} and its submatrices. → Store the row/column offset in a variable k .

```
for(k=0; k<n; k++) {  
    /* Divide by the diagonal element */  
    x[k] = b[k] / L[k+k*ldL];  
  
    /* Subtract L_{*1} x_1 from right-hand side */  
    axpy(n-k-1, -x[k], L+(k+1)+k*ldL, 1,  
        b+(k+1), 1);  
}
```

Forward substitution with BLAS

Goal: Solve $Lx = b$. Using the block notation

$$\begin{pmatrix} \ell_{11} & \\ L_{*1} & L_{**} \end{pmatrix} \begin{pmatrix} x_1 \\ x_* \end{pmatrix} = \begin{pmatrix} b_1 \\ b_* \end{pmatrix},$$



equivalent with $x_1 = b_1 / \ell_{11}$ and $L_{**}x_* = b_* - L_{*1}x_1$.


Recursion: We have to apply the algorithm to the submatrix L_{**} and its submatrices. → Store the row/column offset in a variable k .

```
for(k=0; k<n; k++) {  
    /* Divide by the diagonal element */  
    x[k] = b[k] / L[k+k*ldL];  
  
    /* Subtract L_{*1} x_1 from right-hand side */  
    axpy(n-k-1, -x[k], L+(k+1)+k*ldL, 1,  
        b+(k+1), 1);  
}
```

Forward substitution with BLAS

Goal: Solve $Lx = b$. Using the block notation

$$\begin{pmatrix} \ell_{11} & \\ L_{*1} & L_{**} \end{pmatrix} \begin{pmatrix} x_1 \\ x_* \end{pmatrix} = \begin{pmatrix} b_1 \\ b_* \end{pmatrix},$$



equivalent with $x_1 = b_1 / \ell_{11}$ and $L_{**}x_* = b_* - L_{*1}x_1$.


Recursion: We have to apply the algorithm to the submatrix L_{**} and its submatrices. \rightarrow Store the row/column offset in a variable k .

```
for(k=0; k<n; k++) {  
    /* Divide by the diagonal element */  
    x[k] = b[k] / L[k+k*ldL];  
  
    /* Subtract L_{*1} x_1 from right-hand side */  
    axpy(n-k-1, -x[k], L+(k+1)+k*ldL, 1,  
        b+(k+1), 1);  
}
```

Forward substitution with BLAS

Goal: Solve $Lx = b$. Using the block notation

$$\begin{pmatrix} \ell_{11} & \\ L_{*1} & L_{**} \end{pmatrix} \begin{pmatrix} x_1 \\ x_* \end{pmatrix} = \begin{pmatrix} b_1 \\ b_* \end{pmatrix},$$



equivalent with $x_1 = b_1 / \ell_{11}$ and $L_{**}x_* = b_* - L_{*1}x_1$.

Recursion: We have to apply the algorithm to the submatrix L_{**} and its submatrices. → Store the row/column offset in a variable k.

```
for(k=0; k<n; k++) {  
    /* Divide by the diagonal element */  
    x[k] = b[k] / L[k+k*ldL];  
  
    /* Subtract L_{*1} x_1 from right-hand side */  
    axpy(n-k-1, -x[k], L+(k+1)+k*ldL, 1,  
        b+(k+1), 1);  
}
```

BLAS Level 2

Goal: In order to compute an LU factorization, we have to work with matrices, not vectors.

BLAS Level 2 contains matrix-vector operations, for example

- computing a matrix-vector product, $y \leftarrow y + \alpha Ax$ or $y \leftarrow y + \alpha A^T x$:

```
void gemv(bool trans, int rows, int cols, real alpha,
          const real *A, int ldA,
          const real *x, int incx,
          real *y, int incy);
```

- adding a rank-one product to a matrix, $A \leftarrow A + \alpha xy^T$:


```
void ger(int rows, int cols, real alpha,
          const real *x, int incx,
          const real *y, int incy,
          real *A, int ldA);
```


LU decomposition with BLAS

Goal: Find LU decomposition. Block notation

$$\begin{pmatrix} a_{11} & A_{1*} \\ A_{*1} & A_{**} \end{pmatrix} = \begin{pmatrix} \ell_{11} & \\ L_{*1} & L_{**} \end{pmatrix} \begin{pmatrix} u_{11} & U_{1*} \\ & U_{**} \end{pmatrix}$$

leads to $\ell_{11} = 1$, $u_{11} = a_{11}$, $U_{1*} = A_{1*}$,
 $L_{*1} u_{11} = A_{*1}$, $L_{**} U_{**} = A_{**} - L_{*1} U_{1*}$



In-place algorithm overwrites A directly with L and U .


```
for(k=0; k<n; k++) {  
    scal(n-k-1, 1.0/A[k+k*ldA], A+(k+1)+k*ldA, 1);  
    ger(n-k-1, n-k-1, -1.0,  
        A+(k+1)+ k    *ldA, 1,  
        A+ k    +(k+1)*ldA, ldA,  
        A+(k+1)+(k+1)*ldA, ldA);  
}
```

LU decomposition with BLAS

Goal: Find LU decomposition. Block notation

$$\begin{pmatrix} a_{11} & A_{1*} \\ A_{*1} & A_{**} \end{pmatrix} = \begin{pmatrix} \ell_{11} & \\ L_{*1} & L_{**} \end{pmatrix} \begin{pmatrix} u_{11} & U_{1*} \\ & U_{**} \end{pmatrix}$$

leads to $\ell_{11} = 1$, $u_{11} = a_{11}$, $U_{1*} = A_{1*}$,
 $L_{*1}u_{11} = A_{*1}$, $L_{**}U_{**} = A_{**} - L_{*1}U_{1*}$



In-place algorithm overwrites A directly with L and U .


```
for(k=0; k<n; k++) {  
    scal(n-k-1, 1.0/A[k+k*ldA], A+(k+1)+k*ldA, 1);  
    ger(n-k-1, n-k-1, -1.0,  
        A+(k+1)+ k    *ldA, 1,  
        A+ k    +(k+1)*ldA, ldA,  
        A+(k+1)+(k+1)*ldA, ldA);  
}
```

LU decomposition with BLAS

Goal: Find LU decomposition. Block notation

$$\begin{pmatrix} a_{11} & A_{1*} \\ A_{*1} & A_{**} \end{pmatrix} = \begin{pmatrix} \ell_{11} & \\ L_{*1} & L_{**} \end{pmatrix} \begin{pmatrix} u_{11} & U_{1*} \\ & U_{**} \end{pmatrix}$$

leads to $\ell_{11} = 1$, $u_{11} = a_{11}$, $U_{1*} = A_{1*}$,
 $L_{*1}u_{11} = A_{*1}$, $L_{**}U_{**} = A_{**} - L_{*1}U_{1*}$



In-place algorithm overwrites A directly with L and U .


```
for(k=0; k<n; k++) {  
    scal(n-k-1, 1.0/A[k+k*ldA], A+(k+1)+k*ldA, 1);  
    ger(n-k-1, n-k-1, -1.0,  
        A+(k+1)+ k    *ldA, 1,  
        A+ k    +(k+1)*ldA, ldA,  
        A+(k+1)+(k+1)*ldA, ldA);  
}
```

LU decomposition with BLAS

Goal: Find LU decomposition. Block notation

$$\begin{pmatrix} a_{11} & A_{1*} \\ A_{*1} & A_{**} \end{pmatrix} = \begin{pmatrix} \ell_{11} & \\ L_{*1} & L_{**} \end{pmatrix} \begin{pmatrix} u_{11} & U_{1*} \\ & U_{**} \end{pmatrix}$$

leads to $\ell_{11} = 1$, $u_{11} = a_{11}$, $U_{1*} = A_{1*}$,
 $L_{*1}u_{11} = A_{*1}$, $L_{**}U_{**} = A_{**} - L_{*1}U_{1*}$



In-place algorithm overwrites A directly with L and U .

```
for(k=0; k<n; k++) {  
    scal(n-k-1, 1.0/A[k+k*ldA], A+(k+1)+k*ldA, 1);  
    ger(n-k-1, n-k-1, -1.0,  
        A+(k+1)+ k    *ldA, 1,  
        A+ k    +(k+1)*ldA, ldA,  
        A+(k+1)+(k+1)*ldA, ldA);  
}
```

Summary

Finite differences:

- Replace the domain Ω by a grid Ω_h ,
- replace functions by grid functions, and
- replace differential operators by finite difference operators.

Potential equation approximated by linear system

$$\begin{aligned}-\Delta_h u_h(x) &= f(x) && \text{for all } x \in \Omega_h, \\ u_h(x) &= 0 && \text{for all } x \in \partial\Omega_h.\end{aligned}$$

Direct solvers like the LU factorization can be used to solve this system.

BLAS offers highly optimized routines that can speed up the implementation significantly.