

# Simulation and High-Performance Computing

## Part 14: Explicit Vectorization

Steffen Börm

Christian-Albrechts-Universität zu Kiel

October 6th, 2020

# Intrinsics

**Approach:** Compilers provide special functions that are not implemented in a library, but directly translated into machine code.

**Example:** AVX vectorization

- Definitions can be obtained via `#include <immintrin.h>`
- Data type `__m256` for 256-bit vectors
- Intrinsic functions `_mm256_set1_ps`, `_mm256_loadu_ps`, `_mm256_storeu_ps` to set, load, and store vectors
- Intrinsic functions `_mm256_add_ps`, `_mm256_sub_ps`, `_mm256_mul_ps`, `_mm256_div_ps` for addition, subtraction, multiplication, and division

# Intrinsics

**Approach:** Compilers provide special functions that are not implemented in a library, but directly translated into machine code.

**Example:** AVX vectorization

- Definitions can be obtained via `#include <immintrin.h>`
- Data type `__m256` for 256-bit vectors
- Intrinsic functions `_mm256_set1_ps`, `_mm256_loadu_ps`, `_mm256_storeu_ps` to set, load, and store vectors
- Intrinsic functions `_mm256_add_ps`, `_mm256_sub_ps`, `_mm256_mul_ps`, `_mm256_div_ps` for addition, subtraction, multiplication, and division
- Intrinsic function `_mm256_cmp_ps`, `_mm256_blendv_ps` for comparisons and avoiding branches
- Intrinsic functions `_mm256_and_ps`, `_mm256_or_ps`, `_mm256_xor_ps` for bitwise logical operations

## Example: Reciprocal square root

**Goal:** Evaluate  $x = 1/\sqrt{y}$  efficiently, e.g., for gravitational fields.

**Approach:** The intrinsic function `_mm256_rsqrt_ps` provides us with a rough approximation of  $x$ .

Improve by Newton's iteration for  $f(x) = y - \frac{1}{x^2}$ .

$$x \leftarrow x - \frac{f(x)}{f'(x)} = x - \frac{y - x^{-2}}{-2x^{-3}} = \frac{x}{2} (3 - yx^2)$$

```
const _mm256 c05 = _mm256_set1_ps(0.5f);
```

```
const _mm256 c3 = _mm256_set1_ps(3.0f);
```

```
x = _mm256_rsqrt_ps(y);
```

```
x = _mm256_mul_ps(_mm256_mul_ps(c05, x),
```

```
    _mm256_sub_ps(c3,
```

```
    _mm256_mul_ps(y, _mm256_mul_ps(x, x))));
```

## Example: Reciprocal square root

**Goal:** Evaluate  $x = 1/\sqrt{y}$  efficiently, e.g., for gravitational fields.

**Approach:** The intrinsic function `_mm256_rsqrt_ps` provides us with a rough approximation of  $x$ .

Improve by Newton's iteration for  $f(x) = y - \frac{1}{x^2}$ .

$$x \leftarrow x - \frac{f(x)}{f'(x)} = x - \frac{y - x^{-2}}{-2x^{-3}} = \frac{x}{2} (3 - yx^2)$$

```
const _mm256 c05 = _mm256_set1_ps(0.5f);  
const _mm256 c3 = _mm256_set1_ps(3.0f);  
  
x = _mm256_rsqrt_ps(y);  
x = _mm256_mul_ps(_mm256_mul_ps(c05, x),  
    _mm256_sub_ps(c3,  
    _mm256_mul_ps(y, _mm256_mul_ps(x, x))));
```

**Result:** 0.5 instead of 4.2 seconds, maximal error 2.98<sub>-8</sub>

# Example: Gravitation

Goal: Evaluate gravitational potentials

$$\varphi_i = \sum_{\substack{j=1 \\ j \neq i}}^n \frac{m_j}{\|x_j - x_i\|}.$$

```
d0 = _mm256_sub_ps(_mm256_loadu_ps(y0+j), y0i);
d1 = _mm256_sub_ps(_mm256_loadu_ps(y1+j), y1i);
d2 = _mm256_sub_ps(_mm256_loadu_ps(y2+j), y2i);
norm2 = _mm256_add_ps(
    _mm256_add_ps(_mm256_mul_ps(d0, d0),
        _mm256_mul_ps(d1, d1)),
    _mm256_mul_ps(d2, d2));
force = _mm256_mul_ps(_mm256_loadu_ps(m+j),
    avx_rsqrt(norm2));
sum = _mm256_add_ps(sum, force);
```

## Example: Gravitation

Goal: Evaluate gravitational potentials

$$\varphi_i = \sum_{\substack{j=1 \\ j \neq i}}^n \frac{m_j}{\|x_j - x_i\|}.$$

```
d0 = _mm256_sub_ps(_mm256_loadu_ps(y0+j), y0i);
d1 = _mm256_sub_ps(_mm256_loadu_ps(y1+j), y1i);
d2 = _mm256_sub_ps(_mm256_loadu_ps(y2+j), y2i);
norm2 = _mm256_add_ps(
    _mm256_add_ps(_mm256_mul_ps(d0, d0),
        _mm256_mul_ps(d1, d1)),
    _mm256_mul_ps(d2, d2));
force = _mm256_mul_ps(_mm256_loadu_ps(m+j),
    avx_rsqrt(norm2));
sum = _mm256_add_ps(sum, force);
```

Result: 0.9 seconds instead of 6.9, relative error 1.50<sub>-5</sub>.

## Example: Sine

**Goal:** Evaluate  $y = \sin(x)$  efficiently, e.g., for wave simulations.

**Approach:** Taylor expansion (Horner's method), maximal error  $5.69_{-8}$ .

$$\sin(x) \approx x \left( 1 - x^2 \left( \frac{1}{3!} - x^2 \left( \frac{1}{5!} - x^2 \left( \frac{1}{7!} - x^2 \left( \frac{1}{9!} - \frac{x^2}{11!} \right) \right) \right) \right) \right)$$

```
const float c[] = { 1.0, 1.0/6.0, 1.0/120.0, -.0/5040.0,  
                    1.0/362880.0, 1.0/39916800.0 };  
int j = sizeof(c) / sizeof(float);
```

```
y = _mm256_set1_ps(c[--j]);  
x2 = _mm256_mul_ps(x, x);  
for(; j-->0; )  
    y = _mm256_sub_ps(_mm256_set1_ps(c[j]),  
                      _mm256_mul_ps(x2, y));
```

**Result:** 0.3 instead of 1.4 seconds, maximal error  $1.79_{-7}$



## Example: Sine

**Goal:** Evaluate  $y = \sin(x)$  efficiently, e.g., for wave simulations.

**Approach:** Taylor expansion (Horner's method), maximal error  $5.69_{-8}$ .

$$\sin(x) \approx x \left( 1 - x^2 \left( \frac{1}{3!} - x^2 \left( \frac{1}{5!} - x^2 \left( \frac{1}{7!} - x^2 \left( \frac{1}{9!} - \frac{x^2}{11!} \right) \right) \right) \right) \right)$$

```
const float c[] = { 1.0, 1.0/6.0, 1.0/120.0, -.0/5040.0,  
                   1.0/362880.0, 1.0/39916800.0 };  
int j = sizeof(c) / sizeof(float);
```

```
y = _mm256_set1_ps(c[--j]);  
x2 = _mm256_mul_ps(x, x);  
for(; j-->0; )  
    y = _mm256_sub_ps(_mm256_set1_ps(c[j]),  
                     _mm256_mul_ps(x2, y));
```

**Result:** 0.3 instead of 1.4 seconds, maximal error  $1.79_{-7}$

## Example: Quadratic equation

**Goal:** Stable computation of the solutions of  $x^2 + 2px + q = 0$ .

$$x_1 = \begin{cases} p + \sqrt{p^2 - q} & \text{if } p \geq 0, \\ p - \sqrt{p^2 - q} & \text{otherwise,} \end{cases} \quad x_2 = \frac{q}{x_1}.$$

**Idea:** Copy the sign of  $p$  to the square root.

```
mask = _mm256_castsi256_ps(_mm256_set1_epi32(0x80000000));
sgn = _mm256_and_ps(p, mask);
x1 = _mm256_add_ps(p,
    _mm256_or_ps(sgn,
        _mm256_sqrt_ps(
            _mm256_sub_ps(_mm256_mul_ps(p, p), q))));
x2 = _mm256_div_ps(q, x1);
```

## Example: Quadratic equation

**Goal:** Stable computation of the solutions of  $x^2 + 2px + q = 0$ .

$$x_1 = \begin{cases} p + \sqrt{p^2 - q} & \text{if } p \geq 0, \\ p - \sqrt{p^2 - q} & \text{otherwise,} \end{cases} \quad x_2 = \frac{q}{x_1}.$$

**Idea:** Copy the sign of  $p$  to the square root.

```
mask = _mm256_castsi256_ps(_mm256_set1_epi32(0x80000000));
sgn = _mm256_and_ps(p, mask);
x1 = _mm256_add_ps(p,
    _mm256_or_ps(sgn,
        _mm256_sqrt_ps(
            _mm256_sub_ps(_mm256_mul_ps(p, p), q))));
x2 = _mm256_div_ps(q, x1);
```

**Result:** 0.7 seconds instead of 1.3, maximal error  $2.28_{-7}$ .

# Branches

**Problem:** How can we handle data-dependent branches, if we can only perform operations on entire vectors?

**Idea:** Execute all branches and combine the results.

- `_mm256_cmp_ps` compares all components of two vectors and creates a bitmask:  
If the comparison was successful, the component is set to  $\sim 0$ .  
If it was not successful, the component is set to 0.
- `_mm256_and_ps`, `_mm256_andnot_ps`, `_mm256_or_ps` can be used to combine bitmasks or to merge results according to bitmasks.
- `_mm256_blendv_ps` merges results of two branches according to a bitmask.

## Example: Data-dependent branch

**Goal:** Evaluate  $\sin(\pi x)$  efficiently for  $x \in [0, 1]$ . Our Taylor expansion converges more quickly in  $[0, \frac{1}{2}]$ . Let's use the sine function's symmetry:

$$\sin(\pi x) = \begin{cases} \sin(\pi(1-x)) & \text{if } x \in [\frac{1}{2}, 1], \\ \sin(\pi x) & \text{otherwise.} \end{cases}$$

```
const __m256 c05 = _mm256_set1_ps(0.5f);  
const __m256 c1 = _mm256_set1_ps(1.0f);  
  
msk = _mm256_cmp_ps(x, c05, _CMP_GT_OQ);  
x = _mm256_blendv_ps(x, _mm256_sub_ps(c1, x), msk);  
y = avx_sin_pi(x);
```


The bitmask `msk` for each component is  $\sim 0$  if  $x > \frac{1}{2}$  and 0 otherwise. In the first case, the blend function chooses  $1 - x$ , in the second it chooses  $x$ .

# Latency and throughput

**Latency:** Time from the start to the end of an operation.

**Throughput:** Time from the start of one operation to the start of the next.

**Pipelined processors:** The throughput can be far smaller than the latency.



# Latency and throughput

**Latency:** Time from the start to the end of an operation.

**Throughput:** Time from the start of one operation to the start of the next.

**Pipelined processors:** The throughput can be far smaller than the latency.



**Condition:** The operations must be independent.

# Improving the throughput

**Example:** Taylor expansion by Horner's algorithm.

$$\begin{aligned}y &= a_0 + a_1x + a_2x^2 + a_3x^3 + a_4x^4 + a_5x^5 \\ &= a_0 + x(a_1 + x(a_2 + x(a_3 + x(a_4 + xa_5))))\end{aligned}$$



# Improving the throughput

**Example:** Taylor expansion by Horner's algorithm.

$$\begin{aligned}y &= a_0 + a_1x + a_2x^2 + a_3x^3 + a_4x^4 + a_5x^5 \\&= a_0 + x(a_1 + x(a_2 + x(a_3 + x(a_4 + xa_5)))) \\&= (a_0 + x^2(a_2 + x^2a_4)) + x(a_1 + x^2(a_3 + x^2a_5)).\end{aligned}$$

# Improving the throughput

**Example:** Taylor expansion by Horner's algorithm.

$$\begin{aligned}y &= a_0 + a_1x + a_2x^2 + a_3x^3 + a_4x^4 + a_5x^5 \\&= a_0 + x(a_1 + x(a_2 + x(a_3 + x(a_4 + xa_5)))) \\&= (a_0 + x^2(a_2 + x^2a_4)) + x(a_1 + x^2(a_3 + x^2a_5)).\end{aligned}$$

Mathematically, both formulations are equivalent.

The second formulation allows the left and the right term to be evaluated independently, thus improving the throughput.

# Improving the throughput

**Example:** Taylor expansion by Horner's algorithm.

$$\begin{aligned}y &= a_0 + a_1x + a_2x^2 + a_3x^3 + a_4x^4 + a_5x^5 \\&= a_0 + x(a_1 + x(a_2 + x(a_3 + x(a_4 + xa_5)))) \\&= (a_0 + x^2(a_2 + x^2a_4)) + x(a_1 + x^2(a_3 + x^2a_5)).\end{aligned}$$

Mathematically, both formulations are equivalent.

The second formulation allows the left and the right term to be evaluated independently, thus improving the throughput.

**Problem:** The compiler is not allowed to rearrange floating-point operations if it results in different rounding errors.

→ It is up to us.

# Filling the gaps

**Idea:** We can use the latency gaps of one computation for another one.

**Example:** Sine and cosine computed simultaneously.

```
const float sc[] = { 1.0, 1.0/6.0, 1.0/120.0, 1.0/5040.0,  
                    1.0/362880.0, 1.0/39916800.0 };  
const float cc[] = { 1.0, 1.0/2.0, 1.0/24.0, 720.0,  
                    1.0/40320.0, 1.0/3628800.0 };  
sx = _mm256_set1_ps(sc[5]);  
cx = _mm256_set1_ps(cc[5]);  
x2 = _mm256_mul_ps(x, x);  
for(j=5; j-->0; ) {  
    sx = _mm256_mul_ps(x2, sx);  
    cx = _mm256_mul_ps(x2, cx);  
    sx = _mm256_sub_ps(_mm256_set1_ps(sc[j]), sx);  
    cx = _mm256_sub_ps(_mm256_set1_ps(cc[j]), cx);  
}
```

## Filling the gaps

**Idea:** We can use the latency gaps of one computation for another one.

**Example:** Sine and cosine computed simultaneously.

```
const float sc[] = { 1.0, 1.0/6.0, 1.0/120.0, 1.0/5040.0,  
                    1.0/362880.0, 1.0/39916800.0 };  
const float cc[] = { 1.0, 1.0/2.0, 1.0/24.0, 720.0,  
                    1.0/40320.0, 1.0/3628800.0 };  
sx = _mm256_set1_ps(sc[5]);  
cx = _mm256_set1_ps(cc[5]);  
x2 = _mm256_mul_ps(x, x);  
for(j=5; j-->0; ) {  
    sx = _mm256_mul_ps(x2, sx);  
    cx = _mm256_mul_ps(x2, cx);  
    sx = _mm256_sub_ps(_mm256_set1_ps(sc[j]), sx);  
    cx = _mm256_sub_ps(_mm256_set1_ps(cc[j]), cx);  
}
```

**Result:** Sine takes 0.23 seconds, sine and cosine take 0.34 seconds.

# Prefetching

**Problem:** Read operations from main memory can have very long latencies.

**Idea:** If we know the algorithm well enough, we can provide the processor with hints about what data we will need soon.

```
for(i=0; i+7<n; i+=8) {  
    _mm_prefetch(x+i+64);  
  
    vx = _mm256_loadu_ps(x+i);  
  
    _mm256_stream_ps(y+i, _mm256_mul_ps(vx, vx));  
}
```

- The `_mm_prefetch` intrinsic lets the processor know that we may soon need information from an address.
- The `_mm256_stream_ps` intrinsic lets the processor know that we do not expect to use the written data again soon, so it does not have to be kept in the cache.

# Summary

**Intrinsics** allow us to directly work with the processor's low-level instructions.

Learning how to use intrinsics takes time, but the performance of programs can be significantly improved in situations that are too complicated for the compiler's automatic vectorization.

**Data-dependent Branches** require special treatment, e.g., by constructing bitmasks and using bit manipulation or blending to combine partial results.

**Latencies** during the execution of instructions should be avoided, e.g., by interleaving independent parts of a computation or even completely different computations.

**Prefetching** can help avoid memory latencies, particularly if the access pattern is too irregular for the processor to predict.